

Contents

Reproducible and Collaborative Statistical Data Science:

Overview	2
Administrativia	2
Prerequisites	2
Format and assessment	2
Office hours	3
Graduate Student Instructor	3
Code of conduct; attribution of work	3
Disability accommodations	4
Resources	4
Topics	5
Best uses for Jupyter notebooks	9
Assignments	10
Individual Assignments	10
Assignment 1. Due 9/3, 11:59pm:	10
Assignment 2. Due 9/9, 11:59pm:	10
Reading assignment, finish before class on 9/24	11
Assignment 3. Due 9/30, 11:59pm:	11
Assignment 4. Due 9/30, 11:59pm:	12
Assignment 5. Due 10/7, 11:59pm:	13
Reading assignment, finish before class on 10/17	13
Assignment 6. Due 10/21, 11:59pm:	13
Assignment 7. Due 10/28, 11:59pm:	14
Group Assignments about Ranson (2014) on Climate and Crime.	15
Group Assignment 1. Due 10/21, 11:59pm:	15
Group Assignment 2. Due 11/4, 11:59pm:	15
Group Assignment 3. Due 11/4, 11:59pm:	15
Group Assignment 4. Due 11/4, 11:59pm: (yes, 3 assignments due 11/4)	16
Group Assignment 5. Due 12/2, 11:59pm:	16
Group Assignment 6. Due 12/2, 11:59pm:	17
Group presentation. Due 12/7, 11:59pm:	18
Extra credit (group or individual). Due 12/7, 11:59pm:	18
Collected Reading List:	19

Statistics 159/259: Reproducible and Collaborative Statistical Data Science

Philip B. Stark, Department of Statistics, UC Berkeley

www.stat.berkeley.edu/~stark pbstark@berkeley.edu @philipbstark

Office: 403 Evans Hall. Office hours: Mondays, 12:15-1:15pm

This version: December 3, 2018

Reproducible and Collaborative Statistical Data Science: Overview

This course teaches reproducible and collaborative research techniques through applied statistics, including reproducing published work and re-analyzing the data in that work using other methods-reproducibly.

Examples will be drawn from a variety of fields, including agriculture, health, public policy, and climate.

There will be roughly six small assignments and six larger projects. Much of the work will be collaborative in groups of 4-5. You will be asked to review your own contributions and each others contributions to group projects. There will not be a midterm or final exam, but there will be final presentations of group work.

Administrativia

Prerequisites

- Statistics 133, 134, 135
- Willingness to pick up programming languages and software tools independently (tools used will include Python; Jupyter Notebooks; the Python “scientific stack” of numpy, scipy, matplotlib, and perhaps pandas and scikit; git; GitHub; Travis CI; Docker; LaTeX, Markdown, pandoc)
- Willingness to learn some statistical methodology by reading on one’s own (materials and links will be provided, but not all topics required to do the homework will be covered in lecture)

Format and assessment

- 3 hours of lecture and 2 hours of lab per week (bCourses will have screen-casts of lectures)
 - lectures will focus on theory, applications, and philosophy of science
 - section will focus on computing, software tools, workflow, and collaboration
- approximately 7 individual assignments (70% of grade)
- a group term project, divided into approximately 5 deliverables, plus a final presentation (30% of grade)

Office hours

- Mondays 12:15-1:15, 403 Evans Hall

Graduate Student Instructor

- Mitch Negus, mitchell_negus@b.e
- Office hours 10-12 Tuesdays and Thursdays, 444 Evans Hall

Submitting assignments: Submit written assignments by making a pull request to your private repository within the git organization for the class, <https://github.berkeley.edu/stat-159-259-f18>. Use your CalNet credentials to access your private repository. Create a directory for each assignment labeled with the assignment number, e.g., “Assignment1” for the first assignment.

- Text documents should be written in LaTeX or Markdown. A pdf and the source file should be submitted. Microsoft Word is not acceptable.
- Code and analyses should be in python. All code should have accompanying unit tests. In some cases, Jupyter notebooks will be the appropriate thing to submit; in others (more extensive analyses), a collection of .py files will be more appropriate. For term projects, the “deliverable” will include a repository that includes code, data, analyses, unit tests, and coverage tests.
- Final written projects are due on the last day of final exams, 12/14.

Code of conduct; attribution of work

The high academic standard at the University of California, Berkeley, is reflected in each degree awarded. Every student is expected to maintain this high standard by ensuring that all academic work reflects unique ideas or properly attributes the ideas to the original sources.

These are some basic expectations of students with regards to academic integrity: Any work submitted should be your own individual thoughts, and should not have been submitted for credit in another course unless you have prior written permission to re-use it in this course from this instructor.

All assignments must use “proper attribution,” meaning that you have identified the original source and extent or words or ideas that you reproduce or use in your assignment. This includes drafts and homework assignments! If you are unclear about expectations, ask your instructor.

Do not collaborate or work with other students on assignments or projects unless the instructor gives you permission or instruction to do so.

Disability accommodations

If you need an accommodation for a disability, if you have information you wish to share with the instructor about a medical emergency, or if you need special arrangements if the building needs to be evacuated, please inform the instructor as soon as possible.

If you are not currently listed with DSP (the Disabled Students' Program) and believe you might benefit from their support, please apply online at dsp.berkeley.edu

Resources

- Computing resources
 - We will be using Jupyter notebooks. You can use a hosted notebook at <https://datahub.berkeley.edu/> or install Jupyter on your own device. The datahub.berkeley.edu server will have all the packages you need pre-installed. In contrast, if you use the Anaconda distribution, you will need to install some extra things, such as the permute and cryptorandom packages.
 - We will use the campus github server, github.berkeley.edu
 - The class notes and other materials are available at <https://github.berkeley.edu/pbstark/S159-f18>
 - Assignments should be submitted by pull request to your private repository within the class organization <https://github.berkeley.edu/stat-159-259-f18>
- Git and git workflows
 - Introduction to Git. This is based on the notes we used in this class, but has a fair amount of additional explanation and detail you may find useful through the semester.
 - Immersion course
 - git-scm guide
 - Statlab development git workflow
- Continuous integration
 - Travis CI for beginners
 - Continuous integration with Travis by Simon Scholz
- Scientific Python, Jupyter
 - Lecture notes on scientific python
 - Python for scientific computing by Fernando Perez
 - <https://hplgit.github.io/primer.html/doc/pub/half/book.pdf>
 - Elegant SciPy, Stefan van der Walt. The full book and all the notebooks are available.
 - Getting started with Python for research, a gentle introduction to Python in data-intensive research.
 - An introduction to “Data Science”, a collection of Notebooks by BIDS' Stéfan Van der Walt.

- Effective Computation in Physics, by Kathryn D. Huff; Anthony Scopatz. Notebooks to accompany the book.
- A Whirlwind Tour of Python, by Jake VanderPlas.
- Python for Data Analysis, 2nd Edition, by Wes McKinney, creator of Pandas. Companion Notebooks
- Effective Pandas, a book by Tom Augspurger, core Pandas developer.
- Docker
 - <https://docs.docker.com/get-started/>
 - <https://docker-curriculum.com/>
- LaTeX
 - <https://www.tug.org/twg/mactex/tutorials/ltxprimer-1.0.pdf>
- Markdown
 - <https://daringfireball.net/projects/markdown/syntax>
 - <https://github.com/adam-p/markdown-here/wiki/Markdown-Cheatsheet>
 - <https://www.markdownguide.org/getting-started/>
- Pandoc
 - <https://pandoc.org/getting-started.html>
 - <https://pandoc.org/MANUAL.pdf>
- Miscellaneous computing tutorials
 - Berkeley Statistical Computing Facility tutorials

Topics

- What is reproducibility?
 - Why is reproducibility an issue?
 - Terms in different fields. Computational, experimental, ...
 - “Reproducibility”
 - What is contributing to lack of reproducibility in science?
 - Importance of replication
 - “Virtual witnessing”
- Attempting computational reproducibility in data analysis
 - data
 - * get the data
 - can be hard even if journal/funder requires making data available
 - * figure out what format it’s really in
 - data dictionaries sometimes help
 - proprietary formats common
 - * figure out whether it’s the right data
 - sniff tests
 - consistency tests
 - sleuthing

- * pre-processing, cleaning, etc.
 - analysis
 - * figure out what they claim to have done
 - * figure out what they did
 - usually impossible from just the methods section
 - much harder if the analysis was not scripted
 - might be impossible even with their code
 - * figure out what they should have done
 - * compare
- Sciencing is hard
 - confirmation bias
 - easiest person to fool is yourself
 - misperceptions of probability
 - confidence and accuracy unrelated
 - shiny models and methods
 - broken reward structure
 - ritualization
 - Cargo-cult science and statistics
- Papers/Datasets for the term
 - papers where there seemed to be a chance to get the data
 - topics with social impact: food, health,
 - some paper I think are bogus
 - some papers whose conclusions I like—scrutiny to avoid confirmation bias
- Software engineering
 - revision/version control
 - documentation
 - modularity and abstraction
 - * consistency: APIs, calling signatures, object-oriented coding
 - * separating data, computation, presentation
 - * how general is the problem your approach can solve?
 - what's the right level of abstraction?
 - does it solving it require other broadly useful tools?
 - consider other approaches to subproblems?
 - don't re-invent the wheel...but understand how wheels work
 - unit tests, integration tests, regression tests, coverage tests
 - code review
 - pair programming
 - scripted analyses
 - automation
 - accountability
- Case study: Karp et al., 2015

- access to data
- reproducing the main results from the data
- regression models
 - * assumptions required to perform OLS
 - * assumptions required for OLS to be unbiased
 - * assumptions required to compute SE
 - * assumptions required for $\hat{\beta}/SE$ to have a t-distribution
- interpreting P-values
 - * what's the null hypothesis?
 - * appropriateness of t-tests in regression
 - * p-values from observational data: hypothetical randomness
- permutation tests
 - * group invariances and exchangeability
 - * the Neyman “ticket” model
 - the strong null hypothesis and weak null hypotheses
 - * interference
 - when is non-interference a reasonable assumption?
 - * null hypotheses, tests, and test statistics
 - key restrictions
 - significance versus power
 - specific alternatives and omnibus alternatives
 - p-values versus fixed-level tests
 - * generating random permutations. See Stark 2017
 - generating pseudo-random numbers and pseudo-random integers
 - LCGs, Mersenne Twister, cryptographic PRNGs
 - shuffling algorithms
 - the cryptorandom library
 - problems with R's algorithms for generating random integers and random samples Ottoboni & Stark
 - * confidence bounds for p-values by inverting binomial tests See `Permute/utils binom_conf_interval`
- from reproducibility to replicability, stability, and generalizability
 - * transforming data before regression: “Garden of forking paths.” See Gelman & Loken
 - * sensitivity of conclusions to transformations
 - * sensitivity of conclusions to individual data: “influential observations”
 - * testing before modeling and post-selection inference (POSI)
 - * why reporting everything you tried matters; pre-registration
 - AllTrials.net
 - changing clinical endpoints. Example: PACE trial for CFS/ME
- Sensitivity analysis and sensitivity auditing (guest lectures by Andrea Saltelli and Jeroen van der Sluijs)

- Statistical models and response schedules
 - Response schedules and “physics.” See Freedman SMTP Ch6
 - Linear probability models
 - Logit and probit models
 - Poisson regression
- Goodness of fit tests
- Bayesian and frequentist estimation and inference
 - Foundational issues
 - Interpretation of probability
 - * prior probabilities
 - Example problems
 - * Bounded normal mean
 - * Election auditing
 - Abstract framework
 - Types of uncertainty
 - * Epistemic and aleatory uncertainty
 - * constraints versus priors
 - Bayesian and frequentist measures of uncertainty
 - Duality between minimax and Bayes estimation
- Example: Election audits
 - The auditing challenge and evidence-based elections
 - Public evidence from secret ballots
 - Sequential tests
 - * Wald’s SPRT for Binomial p
 - * Wald’s SPRT for dependent observations
 - Transparency, reproducibility, auditability, and evidence in elections
 - * public observation, public notice
 - observation versus evidence
 - * data disclosure, “commitments”
 - * selecting the seed as a public ritual
 - * source disclosure
 - PRNG
 - mapping from PRNG to sample
 - risk calculations
 - escalation rule and (most importantly) stopping rule
 - * procedural complexity
 - what can a single observer observe?
 - what does the public have to trust to have trust in election outcomes?
 - * examples of disclosed tools
 - tools for ballot-level comparison audits
 - tools for ballot-polling audits
 - SUITE

- RLATool new version
- Contrasting RLAs and Bayesian audits
 - * what question does the audit answer?
 - this election, or a hypothetical population of elections generated from known distribution?
 - * whence the prior?
 - * when are Bayesian audits RLAs?
- Stratified tests and Fisher’s Combining Function
 - why stratify or combine tests?
 - intersection-union tests and union-intersection tests
 - combining functions
 - combinations of independent p -values
 - * chi-square distribution for continuous p -values (under the null)
 - * stochastic dominance by chi-square when the distribution has atoms
 - * dependent tests and lockstep permutations
 - nuisance parameters
 - SUITE

Best uses for Jupyter notebooks

- Jupyter notebooks are a wonderful tool for exploratory data analysis, to present results and to provide a “narrative” analysis: quantitative storytelling, showing the steps of the analysis and explaining the underlying mathematics, science, etc.
- Jupyter notebooks are not an ideal tool to develop a codebase for a project, to house production tools, etc. Jupyter isn’t suited to automated testing or continuous integration (there’s no tool for that in Jupyter, as far as I know). A software development project is generally easier to build and maintain if you separate tests from the code, in different files.
- Jupyter isn’t suitable for packaging/distributing code or for providing tools to be imported into other analyses. For those purposes, you want python files.
- Once you have built the software tools (and tests of those tools) you need for an analysis project, running analyses using those tools in a well documented Jupyter notebook that tells the story of what you did so that others can reproduce it is a good use of the overall tool ecosystem.

Assignments

Individual Assignments

Assignment 1. Due 9/3, 11:59pm:

Getting started reproducing research

1. Look at the data Morabia transcribed from P.C.A. Louis on bloodletting for pneumonia and read Morabia (2006). What do you think of the fact that data from 1828 are available? Reproduce the results below (which Morabia cites); if you cannot reproduce them, say why:
 - 77 patients
 - 2 comparison groups of 41 and 36 patients
 - comparable average age (41 and 38 years, respectively)
 - number of patients bled on the first day, who had passed the age of fifty, was nearly twice as great as that of the patients of the same age, who were bled at a later period
 - duration of disease was an average of 3 days shorter in those who had been bled early compared with those who had been bled late
 - ‘three sevenths’ (i.e., 44 %) of the patients who had been bled early died
 - ‘only one fourth’ (i.e., 25 %) of those bled late died

Is Louis’s work an observational study or an experiment? Do you think it amounts to a “natural experiment”? Why or why not? Give two scientific questions (*statistical hypotheses*) those data might address. What do you think the most important confounding factors would be, for those two hypotheses? What would be the most natural “as-if” randomization to use in analyzing the data to address the hypotheses you formulated, if you were to consider the study to be a natural experiment? What are the controls? Is the experiment blind? Double-blind? (If you are unfamiliar with the notions of confounding, natural experiments, controls, blinding, etc., see the relevant chapters of SticiGui.)

1. Read Karp et al. (2015) and look at the data they provided in Data Dryad. Which figures and tables in the paper could, in principle, be reproduced from the data they provide? Which cannot? Does the methods section describe how they processed the data in adequate detail to reproduce the analyses? If not, what else would you need to know?

Assignment 2. Due 9/9, 11:59pm:

Terminology: Reproducibility, Replicability, Preproducibility, etc.

Read Barba (2018), Buckheit and Donoho (1995), Rokem et al. (2018), Stark (2018). Explain in your own words different senses of the terms “reproducible,” “replicable,” and “repeatable.” In your own words, explain why these concepts are important for science and society. Do you think there’s a reasonable case

for introducing a new term, such as “preproducible” (not necessarily that word, but a new term)? Why or why not? What would you propose as a solution to the problem that different disciplines use “reproducible,” “replicable,” and “repeatable” to mean different things? There’s no length restrictions for this assignment, but I would expect it to take about 2 pages to do a good but concise job.

Reading assignment, finish before class on 9/24

Saltelli et al. (2015), van der Sluijs et al. (2005), van der Sluijs et al. (2008), van der Sluijs (2016)

van der Sluijs and Saltelli will give guest lectures the week of 9/24.

Assignment 3. Due 9/30, 11:59pm:

Sensitivity Analysis, Sensitivity Auditing, and Public Policy

Read and Berk and Freedman (2001), van der Sluijs (2016), Urban (2015)

- Urban reports, “Overall, 7.9% of species are predicted to become extinct from climate change; (95% CIs, 6.2 and 9.8) (Fig 1).”
 - Urban derives his estimate using “Bayesian meta-analysis”
 - * Explain what meta-analysis is, including the assumptions (see Berk and Freedman)
 - * If you can, state the additional assumptions of Bayesian meta-analysis (this might require research)
 - * If you can figure it out from the paper and supplemental materials, state the prior Urban uses
 - Urban points out that there are several general approaches the underlying 131 studies use to estimate the number of species that will go extinct. Sketch how these work:
 - * Species-area relationships
 - * Expert opinion
 - * Species distribution models
 - Recall that the taxonomy of life is Kingdom, Phylum, Class, Order, Family, Genus, Species. There are about 1.9 million known species of eukaryotes (everything but bacteria), and it is estimated that there are 8.7 million in all. There are estimated to be from millions to trillions of species of bacteria (prokaryotes).
 - * Estimate the number of species included in the 131 studies Urban relied on (this might require research: explain how you get your estimate)
 - * Are those species a random sample of all known species? Of all species?
 - * What animal genus has the most species?
 - Do any of the studies Urban relies on examine that genus?

- * What plant genus has the most species?
 - Do any of the studies Urban relies on examine that genus?
- * Do any of the studies Urban relies on consider bacterial species?
- * What families does Urban's estimate consider? (do your best to figure this out—I don't expect you to read all 131 studies)
- What data go into Urban's estimate?
- What, if anything, is random in Urban's estimate?
- Is the estimate of 7.9% of species unbiased? Why or why not?
- Is the range (6.2%, 9.8%) a confidence interval? Why or why not?
- Urban's estimate does not have a timeline: it's "climate change," not "climate change over the next 50y," for instance. How does that make sense?
- List 5 potentially large sources of uncertainty that Urban did not consider or did not address adequately
- On balance, do you think the 7.9% ((6.2%, 9.8%) figures are reliable? Useful? Interpretable?

Assignment 4. Due 9/30, 11:59pm:

Algorithms, unit tests, continuous integration

This assignment concerns the chi-square statistic for the "two-sample problem" for categorical data.

The input is two lists, $x = (x_1, \dots, x_n)$ and $y = (y_1, \dots, y_m)$. Imagine concatenating the lists to get a new list $z = (z_1, \dots, z_N)$ of length $N = n + m$. Let $\{u_k\}_{k=1}^K$ denote the unique values in z and let π_k denote the relative frequency of the value u_k among the elements of z , that is,

$$\pi_k \equiv \frac{\#\{z_j, j = 1, \dots, N : z_j = u_k\}}{N}.$$

Let $E_k \equiv n\pi_k$ and let $O_k \equiv \#\{x_j, j = 1, \dots, n : x_j = u_k\}$ (that is, O_k is the number of elements of x that are equal to u_k). The chi-square statistic for these data for the two-sample problem is

$$\chi^2 = \sum_{k=1}^K \frac{(O_k - E_k)^2}{E_k}.$$

Write three different python functions that (each) take as input x and y and return χ^2 . The functions should use different strategies and/or data structures to calculate χ^2 .

Write unit tests for the functions to ensure that they work correctly for arbitrary input lists x and y .

Assignment 5. Due 10/7, 11:59pm:

Cargo-Cult Statistics and “researcher degrees of freedom”

Read Gelman and Loken (2013) and Silberzahn et al. (2018).

- How many of the co-authors in Silberzahn et al. are in Statistics departments?
- The basic question the teams are supposed to answer is “are soccer referees are more likely to give red cards to dark-skin-toned players than to light-skin-toned players?”
 - Explain in your own words what “more likely” means here.
 - Is there anything random going on? If so, what? What makes it random?
 - What would “equally likely” look like in the real world?
 - Is the question about individual referees, or referees in general?
- For teams 1–6 and 12 (seven teams) in Silberzahn et al., explain the following:
 - Does the analysis implicitly or explicitly involve a model? If it does:
 - * Describe the model
 - * List and explain the assumptions of the model
 - * Assess the evidence that the model is a response schedule
 - * Describe any goodness-of-fit tests the team used to check the adequacy of the model
 - * Explain in words what “OR” means for each of the models.
 - What is OR or what parameter does it estimate?
 - Are the confidence intervals really confidence intervals? Why or why not?

Reading assignment, finish before class on 10/17

In preparation for the guest lecture by Nate Johnson, read Seralini et al., 2014, and skim the comments here.

Assignment 6. Due 10/21, 11:59pm:

This brief assignment (a little computation and a modest amount of thinking) illustrates the problem of selective inference by simulation. Imagine that you are selecting variables to include in a regression model. A common method for selecting which coefficients/variables/parameters to include in a regression or other statistical model is to keep the variable if the estimated coefficient $|\hat{\theta}|$ is “significant,” i.e., if the t-statistic or z-statistic for the estimated coefficient is large.

Having chosen which variables to keep using a method like that, analysts often then report confidence intervals for the coefficients, as if they had not already used the data to decide which variables to keep in the model. This is called “selective inference,” as described in the citations below. Selective inference leads

to problems with statistical reproducibility for many reasons. This assignment highlights one of them: confidence intervals can be far less likely to contain the true value of the parameter when you use the data to select which parameters to report confidence intervals for.

To keep things simple, we will pretend that the standard error of the coefficient is known (and is equal to one), rather than estimated, so instead of using Student's t distribution we can use the standard normal distribution. In this exercise, X plays the role of $\hat{\theta}$. You will simulate $X \sim N(\theta, 1)$; X is then an unbiased estimate of θ .

- For each $\theta \in \{-3, -2.9, \dots, -0.1, 0, 0.1, \dots, 2.9, 3\}$, simulate a draw $X \sim N(\theta, 1)$.
- If the draw is “statistically significant at level 0.05,” i.e., if $|X| \geq 1.96$, construct the usual 95% confidence interval for θ from the draw, i.e., $[X - 1.96, X + 1.96]$. If $|X| < 1.96$, do not make a confidence interval. For each confidence interval, record whether it contains the value of θ used to generate X . Repeat until you have constructed 10,000 confidence intervals for each value of θ .
- Plot the fraction of the 10,000 confidence intervals for θ that contain θ , as a function of θ . (I.e., plot the empirical coverage rate of the confidence intervals.)
- Include unit tests for every function, as usual.
- Explain why the plot looks the way it does.
- Suppose you wanted to create a procedure that had at least 95% coverage probability, no matter what θ is. Sketch how you would have to modify the usual normal confidence interval. I don't expect you to work out the math, just explain heuristically how the confidence interval would have to behave. Would it be symmetric around X ? Would it be longer or shorter than the standard interval? What other changes would you expect? (Hint: consider asymmetric confidence intervals. Extra hint: see Benjamini, Y. and D. Yekutieli, 2005. False Discovery Rate-Adjusted Multiple Confidence Intervals for Selected Parameters, *Journal of the American Statistical Association, Theory and Methods*, 100(469), DOI 10.1198/016214504000001907)

Assignment 7. Due 10/28, 11:59pm:

This assignment is about “the Seralini affair,” discussed in class, including the guest lecture by Nate Johnson. There are two relevant readings in the class literature folder, the republished Seralini et al. (2014) paper itself, and a news item. Seralini et al. (2014) is an open-access publication, but they did not publish their software, and they refused to give it to me (I asked Dr. Seralini for it several times in summer, 2018, in preparation for this course). However, they did make some data available with their republished paper (see “Electronic Supplements” here); the data on tumors and mortality are in the class folder

Data (seralini.xlsx).

- Look at seralini.xlsx (in the class Data folder)
- If you had the software Seralini et al. used, would those data allow you to reproduce figures 4 and 6?

Group Assignments about Ranson (2014) on Climate and Crime.

Every student should make at least 4 commits and at least one pull request for each of these assignments. Every submission should include unit tests for all functionality (using nose or unittest); the coverage of the tests should be at least 99%.

The source we will use for the crime data is here: <https://www.openicpsr.org/openicpsr/project/100707/version/V>
These data have already been cleaned (by Jacob Kaplan, a Ph.D. student in Criminology at U. Pennsylvania); his cleaning scripts are here: https://github.com/jacobkap/crime_data

Group Assignment 1. Due 10/21, 11:59pm:

- Write, document, and test code that takes a collection of values at (lat, long) pairs (intended to represent weather stations) and finds the inverse-distance-weighted average value to another given set of (lat, long) points (intended to represent grid points within a county). This is to replicate Ranson's calculation of the daily temperature in a county. The code should do something sensible if any distance is zero.

Group Assignment 2. Due 11/4, 11:59pm:

- Construct a grid of (lat, long) points within Alameda county separated by approximately 5 miles. The first point should be at (37.905098, -122.272225), near Summit Reservoir.
- Write code to identify all weather stations within x miles of Alameda County
- Identify all weather stations within 10 miles (*not* Ranson's 50 miles) of any of the grid points in Alameda county, and find the weighted average inverse distance from each station to the points in the county grid. Your code for finding the stations should take the distance range as an input parameter (i.e., your code should let you find all stations within 5 miles or 50 miles, too).

Group Assignment 3. Due 11/4, 11:59pm:

- retrieve the weather data for the relevant time periods for stations within 10 miles of any grid point in Alameda County

- identify the stations that meet Ranson’s criteria for inclusion in each year
- calculate the “bias” adjustment for each weather station and for the county
- bin the averaged adjusted temperature data, aggregate it by month using the categories Ranson used

Group Assignment 4. Due 11/4, 11:59pm: (yes, 3 assignments due 11/4)

- split Alameda county into two pieces along the eastern edges of zipcodes 94552 and 94539. Consider all zipcodes within Alameda county that are in or west of either of those zipcodes to be West Alameda and all zipcodes in Alameda that are east of those two zipcodes to be East Alameda. Repeat what you did in group assignments (2) and (3) for East Alameda and West Alameda separately (but using the same grid of points—the original gridpoints in Alameda that are in East Alameda form the grid for East Alameda, and the original gridpoints in Alameda that are in West Alameda form the grid for West Alameda).

Group Assignment 5. Due 12/2, 11:59pm:

Consider weather data from the HCN Berkeley station (ID: USC00040693) and the HCN Livermore station (ID: USC00044997) for the time period covered by Ranson’s work.

- bin the maximum temperature data, separately for the two stations, using the categories Ranson used
- devise and implement a stratified permutation test for the hypothesis that the two cities have “the same weather.” Formulate the hypothesis as a generalized *two-sample problem*, i.e., ask whether differences (between the cities) in the number of days each month in which the maximum temperature is in each bin could reasonably be attributed to chance, if the maximum temperatures had been a single population of numbers randomly split across the two cities.
 - What did you stratify on? Why is that a good choice? Why stratify at all?
 - Combine results across strata using Fisher’s combining function
 - Can you use the chi-square distribution to calibrate the test? Why or why not?
 - Discuss how to take into account simulation uncertainty in estimating the overall P -value
 - **Hint.** `cryptorandom` has a function `getrandbits()` that returns k bits that approximate IID Bernoulli variables with $p = 1/2$. You can think of each bit as the outcome of a coin toss, producing 1 if

the coin landed heads and 0 otherwise. To make your simulation efficient, you will need to write your code so that it vectorizes. In particular, you can use `getrandbits()` to get a coin toss for every day in the overall time period, at one go. If you toss a coin by calling `cryptorandom.random()` and comparing the result to 0.5, your code will run about 255 times slower than if you use `getrandbits()`, since then each “coin toss” involves generating 256 random bits, instead of 1 random bit). If you still can’t get your code to run in a reasonable amount of time, it is OK to use Python’s default PRNG instead of `cryptorandom()`.

- **Hint.** See combining tests
- discuss what your findings mean for Ranson’s approach

Group Assignment 6. Due 12/2, 11:59pm:

- fit the Poisson regression model to the data for all of Alameda County, and for the two pieces of Alameda county separately. Fit the separate estimates simultaneously, including dummy variables for all of Alameda county (treat Alameda County as a whole the way Ranson treated states; East and West Alameda are the two counties in the State of Alameda).
 - **Hint.** If some covariate has the same value in both parts of Alameda in every month (e.g., the number of days with maximum temperature below 10F), do not include it in the model: the corresponding parameter is not identifiable, and the estimation problem will be unstable.
 - **Hint.** `statsmodels` has a GLM function similar to that of R, and has an R-style language for writing formulae
- devise and perform a permutation test to check whether the two pieces of Alameda county are consistent with a single model.
 - explain the particular randomization you are using, its assumptions, and your justification for using it as the null hypothesis
 - try using a cryptographic quality PRNG to simulate random permutations; if you run into computational bottlenecks, it is OK to use Python’s default PRNG instead.
 - **Hint.** One way to ask whether the relationship between weather and crime is different in the two parts of Alameda is to check whether fitting two separate models fits the crime data “surprisingly better” than if the relationship were the same everywhere. The complication is that fitting two models will always fit somewhat better than a single model, because there are more parameters. Consider some combined measure of the fit of the two models to their corresponding data, e.g., RMS error. That’s the test statistic. Now, randomly split the data

into two pieces. For instance, take the two sets of monthly binned weather data and crime data, and toss a fair coin to decide whether East Alameda gets its original data, or the data from West Alameda (West gets whichever East didn't get. The randomization needs to keep the weather and crime for a given month together, so we allocate each entire month of weather and crime data to one location or to the other.) Fit two models to these randomized data, and calculate the RMS error. Repeat many times. If the RMS error of the two models to the randomly assigned data is typically much larger than it is for the original data, that's evidence that the relationship between weather and crime is truly different in the two parts of the county.

Group presentation. Due 12/7, 11:59pm:

Each group should submit a 10-12 minute video presentation covering their work on group assignments 1-6, and their reflections on the Ranson paper. The video presentation, together with the underlying GitHub repository (Python files, Jupyter notebooks, makefile, etc.) are due Friday, 12/7.

The video should be submitted (through GitHub) in a format for which there are free players for Mac OS, Windows, and Linux operating systems (not, e.g., WAV files). Camtasia is highly rated and offers a free trial, but you are welcome to use any application you want to generate the video. Production values are not critical, but please be sure that your visual aids (plots, code, or whatever else you choose to show) are legible on screen, and that the audio is good enough to be understood. Every member of your team should have some screen time or narration time. At a minimum, there should be voice over screenshare, but you are welcome to video yourselves, too.

Include a Makefile in your final submission that runs necessary scripts and notebooks, generates PDFs from notebooks and other analysis documents, and optionally performs data acquisition and cleaning steps. Your Makefile should be able to perform all three of these steps independently or at once, depending on user input. Be sure that target dependencies are properly listed so that changing code or figures will only trigger the build steps necessary to modify your analysis.

Extra credit (group or individual). Due 12/7, 11:59pm:

- devise and implement a conservative goodness-of-fit test for the Poisson regression model. (*Conservative* means the chance of a Type I error is less than or equal to α , and never greater than α .)
 - does it make sense to use a stratified test? if so, what does it make sense to stratify on? how will you combine results across strata?
- test the hypothesis that the data for all of Alameda county arise from the fitted Poisson regression model (note that this is giving the model an advantage, since you are using the fitted parameters); report the p -value

As usual, everything should be documented, tested, scripted, and reproducible.

This assignment is entirely optional. You don't have to submit it at all. If you do choose to do it, it can either be a group effort (if the group wants to take it on) or an individual submission, if you want to go it alone.

Collected Reading List:

Foundations; Statistical Models

1. Feynman, R., 1974. CalTech Commencement Address, <http://calteches.library.caltech.edu/51/2/CargoCult.htm>
2. Freedman, D.A., and D. Lane, 1983. A Nonstochastic Interpretation of Reported Significance Levels, *Journal of Business & Economic Statistics*, 1, 292-298.
3. Freedman, D.A., 1995. Some issues in the foundations of statistics, *Foundations of Science*, 1, 19-39. <https://doi.org/10.1007/BF00208723>
4. Freedman, D.A., 1999. From association to causation: some remarks on the history of statistics, *Statistical Science*, 14(3), 243-258.
5. Freedman, D.A., and R. Berk, 2001. Statistical Assumptions as Empirical Commitments, <http://escholarship.org/uc/item/0zj8s368#page-1> (also in Freedman, D.A., 2010. *Statistical Models and Causal Inference: A dialog with the Social Sciences*, Cambridge University Press. D. Collier, J. Sekhon, P.B. Stark, eds.)
6. Freedman, D.A., 2008. On types of scientific inquiry: the role of qualitative reasoning, *The Oxford Handbook of Political Methodology*, Box-Steffensmeier, J.M., H.E. Brady, and D. Collier (eds), Oxford University Press, Oxford. DOI: 10.1093/oxfordhb/9780199286546.003.0012. Preprint
7. Freedman, D.A., 2009. *Statistical Models: Theory and Practice*, 2nd edition, Cambridge University Press.
8. Freedman, D.A., R. Pisani, and R. Purves, 2007. *Statistics*, 4th edition, W.W. Norton, New York.
9. Klemes, V., 1989. The Improbable Probabilities of Extreme Floods and Droughts, in O. Starosolsky and O.M. Meldev (eds), *Hydrology and Disasters*, James and James, London, 43-51. https://www.itia.ntua.gr/en/getfile/1107/1/documents/1997_ImprobProbabilities_OCR.pdf
10. LeCam, L., 1977. Note on metastatistics or 'An essay toward stating a problem in the doctrine of chances,' *Synthese*, 36, 133-160.
11. Stark, P.B., 1997. SticiGui

12. Stark, P.B., 2016a. Pay no attention to the model behind the curtain
13. Stark, P.B., 2016b. The value of P-values, *The American Statistician*, 70, DOI:10.1080/00031305.2016.1154108
14. Stark, P.B., 2017. Mathematical Foundations, Inequalities, Statistical models, Introduction to permutation tests, Rabbits and Cargo-Cult Statistics, Generating pseudo-random samples and permutations Combining tests
15. Stark, P.B., and A. Saltelli, 2018. Cargo-cult Statistics and Scientific Crisis, *Significance*, 15(4), 40–43. <https://www.significancemagazine.com/593>

Statistical methodology

1. Freedman, D.A., 2008. Randomization does not justify logistic regression, *Statistical Science*, 23 237–249. DOI: 10.1214/08-STS262 <https://arxiv.org/pdf/0808.3914.pdf>
2. Hastie, T., R. Tibshirani, and J. Friedman, 2009. *Elements of Statistical Learning: Data Mining, Inference, and Prediction. Second Edition*, Springer-Verlag, NY. <https://web.stanford.edu/~hastie/Papers/ESLII.pdf>
3. McCullagh, P. and J.A. Nelder, 1983. *Generalized Linear Models*, 2nd edition, Chapman & Hall, NY.

Evidence, Models, and Public Policy

1. Saltelli, A., P.B. Stark, W. Becker, and P. Stano, 2015. Climate Models as Economic Guides: Scientific Challenge or Quixotic Quest?, *Issues in Science and Technology*, Spring 2015. Reprint: <http://www.stat.berkeley.edu/~stark/Preprints/saltelliEtal15.pdf>
2. van der Sluijs, J.P., J.S. Risbey, and J.R. Ravetz, 2005. Uncertainty Assessment of Voc Emissions From Paint in the Netherlands Using the NUSAP System, *Environmental Monitoring and Assessment*, 105, 229–259. doi:10.1007/s10661-005-3697-7
3. van der Sluijs, J.P., A.C. Petersen, P.H.M. Janssen, J.S. Risbey, and J.R. Ravetz, 2008. Exploring the quality of evidence for complex and contested policy decisions, *Environmental Research Letters*, 3, doi:10.1088/1748-9326/3/2/024008
4. van der Sluijs, J.P., 2016. Numbers Running Wild, Chapter 5 in *The Rightful Place of Science: Science on the Verge*, A. Benessia, S. Funtowicz, M. Giampietro, Â.G. Pereira, J. Ravetz, A. Saltelli, R. Strand, J.P. van der Sluijs, eds., Consortium for Science, Policy & Outcomes, AZ & DC. http://www.andreasaltelli.eu/file/repository/Science_on_the_Verge_FINAL_.pdf

Foundations: Computation, Optimization

1. Goldberg, D., 1991. What every computer scientist should know about floating-point arithmetic, *ACM Computing Surveys*, 23, 5–48.

2. Lawson, C.L. and R.J. Hanson, 1974. *Solving Least Squares Problems*, Prentice-Hall, NJ.
3. Ottoboni, K. and P.B. Stark, 2018. Random problems with R, ArXiv, <https://arxiv.org/abs/1809.06520>. also see <https://stat.ethz.ch/pipermail/r-devel/2018-September/076817.html>

Agriculture, Ecology, and Health

1. Fagan, J., T. Traavik, and T. Bøhn, 2015. The Seralini affair: degeneration of Science to Re-Science?, *Environmental Sciences Europe*, 27:19, DOI 10.1186/s12302-015-0049-2
2. Karp, D.S., S. Gennet, C. Kilonzo, M. Partyka, N. Chaumont, E.R. Atwill, and C. Kremen, 2015. Comanaging fresh produce for nature conservation and food safety, *PNAS*, 112 (35) 11126-11131. <https://doi.org/10.1073/pnas.1508435112>
3. LeCanne, C.E., and J.G. Lundgren, 2018. Regenerative agriculture: merging farming and natural resource conservation profitably, *PeerJ* 6:e4428 <https://doi.org/10.7717/peerj.4428>
4. Morabia, A., 2006. Pierre-Charles-Alexandre Louis and the evaluation of bloodletting, *J. Roy. Soc. Medicine*, 99, 158–160. <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC1383766/pdf/0158.pdf>
5. Seralini, G.-E., E. Clair, R. Mesnage, S. Gress, N. Defarge, M. Malatesta, D. Hennequin, and J. Spiroux de Vendômois, 2014. Republished study: long-term toxicity of a Roundup herbicide and a Roundup-tolerant genetically modified maize, *Environmental Sciences Europe*, 26:14, <http://www.enveurope.com/content/26/1/14>

Pedestrians and Race

1. Coughenour, C., S. Clark, A. Singh, E. Claw, J. Abelar, and J. Huebner, 2017. Examining racial bias as a potential factor in pedestrian crashes, *Accident Analysis and Prevention*, 98, 96-100. <http://dx.doi.org/10.1016/j.aap.2016.09.031>
2. Goddard, T., K.B. Kahn, and A. Adkins, 2015. Racial Bias in Driver Yielding Behavior at Crosswalks, *Transportation Research Part F: Traffic Psychology and Behaviour*, 33, 1-6. <http://dx.doi.org/10.1016/j.trf.2015.06.002>

Earthquake probabilities

1. USGS 2008 Bay Area Earthquake Probabilities. <http://earthquake.usgs.gov/regional/nca/ucrf/>
2. Cornell, C.A., 1968. Engineering seismic risk analysis, *Bull. Seism. Soc. Am*, 58, 1583–1606.

3. Mulargia, F., P.B. Stark, and R.J. Geller, 2017. Why is probabilistic seismic hazard analysis (PSHA) still used? *Physics of the Earth and Planetary Interiors*, 264, 63-75. <https://doi.org/10.1016/j.pepi.2016.12.002>
4. Stark, P.B. and D.A. Freedman, 2003. What is the Chance of an Earthquake? in *Earthquake Science and Seismic Risk Reduction*, F. Mulargia and R.J. Geller, eds., NATO Science Series IV: Earth and Environmental Sciences, v. 32, Kluwer, Dordrecht, The Netherlands, 201–213. Preprint: <http://www.stat.berkeley.edu/~stark/Preprints/611.pdf>

Impact of Climate Change

1. Houser, T., R. Kopp, S. Hsiang, M. Delgado, A. Jina, K. Larsen, M. Mastrandrea, S. Mohan, R. Muir-Wood, D.J. Rasmussen, J. Rising, and P. Wilson, 2014. The American Climate Prospectus: Economic Risks in the United States, 2014. <http://rhg.com/reports/climate-prospectus>
2. Hsiang, S., R. Kopp, A. Jina, J. Rising, M. Delgado, S. Mohan, D.J. Rasmussen, R. Muir-Wood, P. Wilson, M. Oppenheimer, K. Larsen, and T. Houser, 2017. Estimating economic damage from climate change in the United States, *Science*, 356, 1362-1369 DOI: 10.1126/science.aal4369
3. Ranson, M., 2014. Crime, weather, and climate change, *Journal of Environmental Economics and Management*, 67(3), 274-302. <https://doi.org/10.1016/j.jeem.2013.11.008>
4. Urban, M.C., 2015. Accelerating extinction risk from climate change, *Science*, 348, Issue 6234, 571–573, DOI: 10.1126/science.aaa4984, <http://science.sciencemag.org/content/348/6234/571.full>

Reproducibility and Scientific Method

1. Ball, P., 2018. High-profile journals put to reproducibility test, *Nature*, 27 August, <https://www.nature.com/articles/d41586-018-06075-z>, doi: 10.1038/d41586-018-06075-z
2. Barba, L., 2016. The hard road to reproducibility, *Science*, 354, 142. doi 10.1126/science.354.6308.142
3. Barba, L., 2016. Reproducibility Syllabus, <http://lorenabarba.com/blog/barbagroup-reproducibility-syllabus/>
4. Barba, L., 2018. Terminologies for Reproducible Research, <https://arxiv.org/abs/1802.03311>
5. J.B. Buckheit and D.L. Donoho, 1995. Wavelab and Reproducible Research, https://statweb.stanford.edu/~wavelab/Wavelab_850/wavelab.pdf
6. Implementing Reproducible Research, edited by V. Stodden, F. Leisch and R. Peng
7. The Practice of Reproducible Research: Case Studies and Lessons from the Data-Intensive Sciences Kitizes, Turek and Deniz, eds.

8. Reproducibility: a Primer on Semantics and Implications for Research
9. Mesnard, O., and L.A. Barba, 2017. Reproducible and replicable CFD: it's harder than you think, *IEEE/AIP Computing in Science and Engineering*, 19(4):44–55. 10.1109/MCSE.2017.3151254 Preprint: arXiv:1605.04339
10. Rokem, A., B. Marwick, and V. Staneva, 2018. Assessing Reproducibility, in *The Practice of Reproducible Research: Case Studies and Lessons from the Data-Intensive Sciences*, University of California Press. <https://www.practicereproducibleresearch.org/core-chapters/2-assessment.html>
11. Schapin, and Schaffer, 1985. *Leviathan and the Air-Pump: Hobbes, Boyle, and the Experimental Life*, Princeton University Press, NJ ISBN 0-691-08393-2
12. Stark, P.B., 2018. No Reproducibility Without Preproducibility, *Nature*, 557, 613. <https://www.nature.com/magazine-assets/d41586-018-05256-0/d41586-018-05256-0.pdf> doi: 10.1038/d41586-018-05256-0
13. Stark, P.B., 2017. Preface to *The Practice of Reproducible Research*, J. Kitzes, D. Turek, and F. Deniz, eds., University of California Press, Berkeley
14. Teytelman, L., 2018. No more excuses for non-reproducible methods, *Nature*, 560, 411. <https://www.nature.com/articles/d41586-018-06008-w>, doi: 10.1038/d41586-018-06008-w

Weaponizing Reproducibility

1. Hakim, D. and E. Lipton, 2018. Pesticide Studies Won E.P.A.'s Trust, Until Trump's Team Scorned 'Secret Science', *New York Times*, 26 August. <https://www.nytimes.com/2018/08/24/business/epa-pesticides-studies-epidemiology.html>